

A Machine Reads the Hopper Ledgers: Constrained Transcription, Mechanical TEI Export, and Checkability as the Editorial Condition

Michel Hua

Commutator, Paris

michel@commutator.io

Draft of 20 September 2026, prepared for the *Journal of the Text Encoding Initiative*

Abstract

This article reports an experiment in the transcription and encoding of an artist's record books by a large language model under editorial constraints, using as its corpus the six ledgers of Edward and Josephine Hopper at the Whitney Museum of American Art (504 digitised sheets, 1907–1967) and four of Josephine Hopper's notebooks. The model (Claude Opus 5 and Claude Fable 5.1) transcribed the sheets from the museum's published photographs into a deliberately restricted LaTeX whose macros correspond one to one with elements of a TEI P5 customisation; a mechanical conversion produces TEI files that are validated against the customisation in continuous integration; scripts derive from the transcriptions, under published rules, an account of prices, materials, formats and dated days; and an interface built for the purpose places the reading beside the photograph so that any sentence can be checked against the hand it was read from. The article describes the corpus, the transcription rules, the encoding and its schema, the derived data and the interface, and then the limits: the reading is a first pass unchecked by a person, an alignment pass on one volume reports seventy-four leaves with at least one line the transcription lacks, and the language model's pass did not record the ink colour that carries the cash book's meaning, which an image pass recovered only in part. The contribution is not a finding about Hopper but an editorial condition: a machine reading fully cited, fully derived and fully exposed to correction. The use of generative AI in the transcription and in the drafting is declared.

Keywords. artist's ledgers; large language models; transcription; critical apparatus; hand attribution; Edward Hopper; reproducibility.

1 Introduction

Edward Hopper (1882–1967) recorded his finished work in a set of ledgers that his wife, the painter Josephine Nivison Hopper (1883–1968), kept from their marriage in 1924 until his death. Six volumes survive at the Whitney Museum of American Art, accessioned as 96.208 to 96.213, and the museum has digitised and published every sheet with a catalogue descriptor (Whitney Museum of American Art, n.d.). The volumes are heterogeneous. Three are work books in which Edward Hopper drew a small ink record of each finished picture with its title and size, and Josephine Hopper wrote everything else: the paint used, a description, the buyer, the price, the dealer's commission, the exhibitions and their verdicts, and the date the cheque cleared. One is a pocket cash book kept without a gap from 15 November 1913 to

23 March 1967, in a hand that predates the marriage by eleven years and that the present edition does not attribute. One is a book of drawings and etching stock-takes, and one is a cash book indexed by dealer. Beside them, the Whitney’s Sanborn Hopper Archive holds some ninety notebooks of Josephine Hopper’s, four of which have been digitised: a diary of the marriage as it bore on her own painting, two working files from the 1947 fight against New York University’s attempt to evict the tenants of 3 Washington Square North, and a two-ring memorandum book of lists.

The ledgers have been used before. Gail Levin’s catalogue raisonné was written from the record books and her biography from the diaries (Levin 1995a, 1995b); Deborah Lyons published the record books in 1997 (Lyons 1997); Madeleine Larson’s typescript of the Provincetown diaries exists in five parts (Larson n.d.). What has not existed is a transcription of the digitised sheets that can be placed beside the photograph of each sheet and checked against it, with an apparatus that states what was read, what was doubted, what could not be read and whose hand wrote it. This article reports an experiment in producing such a transcription with a large language model as the first transcriber, under constraints designed to make its output checkable rather than trustworthy.

The objective is methodological, and the object of the article is the method and its apparatus: the transcription is their demonstration, and no claim about Hopper is made here that the demonstration would need to be correct for. The Hopper corpus was chosen because it is the right size for a test, because the museum serves its photographs openly at a resolution at which Josephine Hopper’s figures are legible, and because a ledger is a document in which errors are measurable: a price either adds up or it does not. The article therefore presents the corpus (Section 3), the transcription method and its rules (Section 4), the TEI encoding and the customisation that validates it (Section 5), the interface and the derived data that make the reading reproducible (Section 6), a discussion of what the method permits and what it does not (Section 7), and the same method run against three sibling corpora — two of mathematics, one a filmed lecture — which is where the editorial condition either generalises or fails (Section 8).

Four words are kept apart throughout. A reading is *reproducible* when the process that made it can be rerun; *auditable* when a reader can find the evidence for each sentence; *reliable* when like cases are treated alike; and *valid* when it is right. This edition claims the first three and not the fourth: nothing below makes a wrong reading right, only findable.

2 Related work

Four bodies of work bear on this experiment: account records in TEI, editions of working notebooks, the transcription of handwriting by large language models, and the evaluation of such models as producers of TEI. A fifth, the attribution of hands, is taken up in Section 7.

The TEI encoding of account books has a small but explicit literature. Tomasek and Bauman (2013), writing in this journal, proposed a stand-off model of financial records, the «transactionography», in which transfers between parties are represented as data linked by pointers to the transcribed document rather than embedded in its markup, so that a transaction can be encoded once even where the document records it in several places. That distinction between the document and the transaction is the one the present edition adopts, with a different division of labour: the transcription encodes the document and nothing but the document, and every transaction is recovered from it by a script whose rules are published and whose

output is regenerated on each build (Section 6). The present corpus is also unlike a double-entry book in that its writer kept no consistent form; the reader that recovers transactions from Book IV has to change its rule four times across the volume's fifty-four years, and Section 6.2 describes why.

Working notebooks have a literature of their own, and it is the nearest to this corpus. Scholger (2020), in this journal, encodes the thirty-five notebooks of the conceptual artist Hartmut Skerbisch — some 2,100 inscribed pages, about 450 of them carrying drawings — and argues that a plain `<graphic>` loses what the drawing states: her model separates the graphical component, the textual functions attached to it, and an interpretative layer addressed through `@ana`, so that the stages by which a sketch became a work can be followed across the corpus. The present edition does markedly less with its drawings: Edward Hopper's 205 record sketches are memoranda of finished pictures and not stages of them, so each is marked as a figure, attributed, and left undescribed, since a description of a drawing one pane away competes with looking at it. Scholger's model is what this edition would need were the sketches read as drawings.

The transcription of handwriting by large language models has moved quickly. Humphries and colleagues (2025) report character error rates of about six to seven per cent for general-purpose language models on eighteenth- and nineteenth-century English manuscripts, an improvement on specialised handwritten text recognition; Crosilla, Klic and Colavizza (2025) benchmark several models on the same task and discuss when the models' linguistic priors help and when they hinder. The present experiment differs in its object and its measure: not running prose but ruled tables of figures, names and abbreviations, where a language model's prior is a liability, since a plausible name is not a correct one and a rounded figure is a false transaction; and not an error rate against a ground truth, which does not exist for these sheets, but the reconciliation of the transcribed figures with the sums the writer pencilled at the foot of the leaves.

The evaluation of language models as producers of TEI is more recent. Strutz (2026) proposes a multi-dimensional framework for assessing model-generated TEI, from syntactic validity to semantic accuracy, on historical correspondence. The present edition does not ask the model to produce TEI at all: it writes a restricted language of sixteen macros from which the TEI is derived mechanically, so that the encoding's validity is a property of the converter and a difference between the two files is a bug rather than a judgement. Section 5 sets out the reasoning and Section 7 its cost.

3 The corpus

3.1 The six ledgers

The Whitney's six volumes are listed in Table 1. Three characteristics of the corpus shaped the method.

First, three numbering systems overlap on every sheet and only one of them is an address. The Hoppers wrote leaf numbers on the paper; the museum's digitisation counts photographs, covers and versos included, so that it agrees with the leaf numbers nowhere; and the museum's ResourceSpace holds a resource reference that is in no order but that names exactly one image. «Leaf 58» names six different things across the six volumes and three photographs within Book I alone. A citation from the edition therefore prints both the leaf, which is what a reader of the photograph can see, and the reference, which is the only thing the museum can be asked

Table 1: The six ledgers as the Whitney Museum identifies them, with the number of digitised sheets and the state of the transcription on 13 September 2026.

Volume	Accession	Dates	Contents	Sheets
Book I	96.208	1913–1963	Etchings, one plate to a leaf; running lists from leaf 45	117
Book II	96.209	1907–1962	Paintings, one to an opening, with record sketches	72
Book III	96.210	1924–1967	Late work; lists at the front; memoranda	90
Book IV	96.211	1913–1967	Pocket cash book, continuous	161
Book V	96.212	1953–1963	Drawings; etching stock-takes	16
Dealers/Etchings	96.213	1921–1951	Prints by dealer; leaves numbered from 51	48

about.

Second, the volumes were written by two people and annotated by more. The inside front cover of Book I states the division: «Recorded by Jo N Hopper (Mrs. Edward Hopper) at time each work is finished, or before it leaves studio. Drawings in the 3 books done by Edward Hopper.» The size in the title line of each work is the one figure Edward Hopper wrote; everything else in the work books is hers. Book IV is in a hand the edition cannot place. A transcription that flattened the hands would destroy the archive’s chief documentary value, and the attribution of hands, which most transcription apparatus treats as optional, is accordingly the one element this method treats as mandatory.

Third, the colour of ink carries information. In Book IV, from leaf 7 onward, charges are written in black, receipts in red, and the sums the writer made herself in pencil; on the etching leaves of Book I, the writer states on the index leaf that ink marks the date a cheque was received and pencil the date a picture was sold. A transcription in one typeface loses this unless it records the colour, and a language model reading a photograph does not reliably distinguish it (Section 7).

3.2 The four notebooks

The four notebooks digitised by the Whitney are listed in Table 2. They are not of one kind. *Garrulities + Grouch* is a diary of twenty-two entries, eighteen of them dated, between 1927 and 1958, with the dates in the margin and the entries not in chronological order. *Re: 3 Wash Sq* is a working file of drafts and a chronicle, in which the drafts were written from the front of the pad and the chronicle from the back, so that the two meet in the middle and the chronicle runs backwards through the photographs. *Battle of Wash Sq.* is a memorandum book of eight sheets. The black two-ring notebook is a book of lists, whose dated content runs from 1937 to 1964 and whose subject changes every page or two. Their encoding required one element the ledgers do not: a dated entry, with the writer’s date and the transcriber’s assigned date kept apart.

Table 2: The four notebooks, Sanborn Hopper Archive, Series IV, Subseries A, with the Whitney’s dates and the dates of their contents as read.

Notebook	Whitney’s date	Contents as read	Sheets
Garrulities + Grouch	1924–1950	1927–1958	42
Re: 3 Wash Sq	1947	January–June 1947	51
Battle of Wash Sq. (EJHA.1597)	1947	1947–1955	8
Black two-ring notebook	c. 1952–1961	1937–1964	68

3.3 What was read, and what was not

Table 3 keeps the article’s counts apart: sheets digitised, sheets transcribed, and the files and passes the reading took. The six sheets without a transcription are five back covers or chip-board and one notebook’s front cover, none carrying writing, recorded as not transcribed rather than as empty. A batch is a file; a pass is one sitting of the model on it. No batch is declared checked.

Table 3: The corpus as counted, on 20 September 2026: sheets digitised by the museum, sheets transcribed, and the files and passes the reading took.

	Digitised	Transcribed	Files	Passes
Six ledgers	504	499	44	46
Four notebooks	169	168	4	16
Total	673	667	48	62

3.4 Provenance and access

The three work books and Book V were the gift of Lloyd Goodrich, who organised the Whitney’s Hopper retrospectives of 1950 and 1964; Book IV and the dealers’ book were purchases. The notebooks came with the Sanborn Hopper Archive in 2017. Twenty-two further diaries of 1933 to 1956 are at the Provincetown Art Association and Museum, which publishes Madeleine Larson’s typescript of them; those are not transcribed here and are used only through the typescript, cited by part and page.

The edition stores no photograph. The Whitney’s download endpoint answers cross-origin requests and honours byte ranges; the interface points an image element at the museum’s file, fetched as it is looked at, and the repository refuses image files. The inventory of sheets was harvested by a person in a browser, because the museum’s listing pages sit behind a bot-check that the edition does not circumvent, and the harvest is committed as text so that it can be reviewed in a diff.

4 The transcription method

4.1 The schema, and its values here

Figure 1 gives the whole arrangement at a glance, and it is stated generally because it was not designed for this corpus: the same five stations and the same loop run the three sibling editions of Section 8, and what distinguishes the four is only the values the stations take. A *witness* is what the holding institution serves, addressed by the holder’s own reference, and it is never committed. A *reading* is the one file that is written — by the model first, by a person thereafter — in a language whose every construct is declared, in passes whose header records the model and the date. The *derivations* are everything else the edition publishes, computed from the reading by scripts and regenerated on every build. The *verifications* are whatever a machine can check of them: a schema, an arithmetic, the agreement of two readers. The *address* is the reading placed beside its witness, each sentence cited by the holder’s reference and the edition’s dated pass. The loop is the person, who compares the reading with the witness, corrects the reading, and declares that the comparison was made.




Figure 1: The schema the four editions of Section 8 share, and its values in this one. Five stations and one loop: the model turns a witness into a reading, in passes whose header names the model and the date; scripts derive everything else from the reading and regenerate it on every build; checkers fail the build on what a machine can check; the site places the reading beside the witness and cites both by reference. The dashed loop is the person, who compares the reading with the witness and corrects the reading, which alone changes anything downstream. Four invariants hold at every station: the witness is never committed; only the reading is written, everything else is derived and compared; every assertion carries an address into the witness; and the person’s verdict is declared, not observed. The grey row gives this edition’s values; the siblings’ are in Table 8.

The four invariants the caption states are each enforced by a script or a rule in every edition, and they leave a corpus few things to set: what the witness is and how it is addressed; whether the model is to read it, to read and restate it, or to index a reading already made; the language and the size of a pass; which derivations are worth publishing; which verifications the material admits; and what the person’s check consists of. Table 8 gives the four editions’ values side by side.

Here the witness is the Whitney’s photograph of a sheet, addressed by its ResourceSpace reference and fetched by the browser as it is looked at; the model reads it, twelve sheets to a pass, into a LaTeX of seventeen macros; the derivations are the TEI, the typeset and rendered readings, and four censuses; the verifications are the customised schema, the Lean certificates and the regeneration on every push; and the address is the site, which cites a sheet, a leaf and a pass separately. The person’s file is empty. The rest of this section takes the stations in turn.

4.2 The unit of work and the model

A batch is twelve sheets, read in one fresh model context, and written as a single LaTeX file whose header records the model, the date, and what the batch found difficult. The skill that performs the reading accepts three models, Claude Fable 5.1, Claude Fable 5 and Claude Opus 5, and refuses any other, on the ground that a pass made on a model nobody chose produces a file whose provenance is an accident. Every file records the model that read its sheets, so that provenance is a fact about the file and not about the session. Between 5 and 13 September 2026, Opus 5 read most of the ledgers and the first sittings of the notebooks; Fable 5.1 read the last sittings of the black notebook and the Battle of Wash Sq. Where the rule of one context per batch was not kept, the header says so: batches 3 and 4 of Book I shared a context.

4.3 The rules

The governing rule is stated in the edition’s method and is specific to ledgers: a guessed figure is a sale that did not happen. An invented word in a letter is a bad reading; an invented price in a ledger is a transaction at a price nobody paid, to a buyer who never bought, and it will be

cited, because a table looks like data in a way that prose does not. From that rule follow the others.

Nothing illegible is guessed. A passage that cannot be read is marked illegible, and the mark is never replaced by a plausible reading. A reading offered with doubt is marked doubtful. Nothing is expanded or normalised: «Inv.» stays «Inv.», «30 - 1/3» stays as written, and a name spelt four ways in the archive stays spelt four ways, because a corrected name is a claim about who bought a picture that the leaf did not make. A key to the abbreviations and notations is published separately, and each of its entries must match a transcribed leaf or the build fails. What the writer struck out is kept and marked struck, since a crossing-out is often the more informative half of an entry. Marginal additions are transcribed as marginal, with their position. The transcriber's own observations are confined to a note element that is typographically and structurally distinct from the writer's text, so that a count over the corpus can exclude it; the materials census, for example, reads only the writer's words, because the transcriber's notes discuss materials constantly and a count that read them would report a rare colourman as common because somebody remarked on its rarity.

Hands are attributed to Josephine Hopper, to Edward Hopper, to a later writer, or to nobody. *Unidentified* is a permitted answer and is, the method notes, the right one far more often than confidence suggests. The drawn lettering of the covers is unattributed; the whole of Book IV is unattributed; a pencil annotation on Book II leaf 46 in a heavier hand is attributed to nobody. The edition's own attributions rest on a document and not on letterforms (Section 7.2); where a model attributes a passage on the letterforms alone, its note says so.

There is one edition and no second register. The Grothendieck edition this method derives from (Section 8) carries a modernised text beside the diplomatic one; a ledger needs none, since its English is plain and its columns are already a table, and a second text would be a second artefact to keep in step.

4.4 What the model wrote

The model writes a restricted LaTeX defined in a preamble of seventeen macros (Table 4). The restriction is deliberate. A converter that accepted arbitrary LaTeX would silently mangle what it did not understand, and a silently mangled ledger is a table of numbers that look right; the renderer accepts only the subset, with ordinary line breaks and spacing, and fails loudly outside it, naming the file, the line and the construct. It also refuses a resource reference that is not in the harvested inventory and a file that omits the rights notice.

The choice of LaTeX rather than XML as the model's output language was practical and is discussed in Section 7. Its consequence for encoding is that the model never writes an element name, an attribute or a namespace, and cannot therefore produce ill-formed or invalid markup itself — the converter can, which is what the validation of Section 5.5 exists to catch; what the model can do is misapply a macro, and the converter's checks and the schema's closed value lists are designed to catch the misapplications that have a structural signature.

4.5 A transcribed leaf

The following is the writer's description of Nighthawks, from Book II leaf 95, as the model read it, given here in the TEI that the converter produces from the transcription (the sheet is ref 17982). The struck word, the marginal additions with their positions, the hand and the transcriber's note are all present, and the museum's reference is carried in the page break.

Table 4: The transcription macros and the TEI elements the converter maps them to.

Macro	Meaning	TEI
<code>\sheet{ref}{leaf}</code>	A sheet begins: the museum's reference, the leaf written on the paper	pb with @n, @facs, @xml:id
<code>\hand{who}{...}</code>	Whose hand: jo, edward, later, unidentified	seg with @hand, and handNote
<code>\ill{}</code>	Illegible; never guessed	gap reason="illegible"
<code>\uncertain{...}</code>	A reading offered and flagged	unclear
<code>\struck{...}</code>	Struck out in the book	del
<code>\supplied{...}</code>	What the transcriber supplies and the leaf does not carry: a letter completing a word, an implied sign, a column heading the writer ruled but never labelled	supplied with @resp
<code>\add{...}</code>	The writer's own insertion: a price revised above the line, a buyer written over a struck one	add
<code>\note{...}</code>	The transcriber's note	note type="editorial"
<code>\marginal{...}</code>	The writer's marginal note	note place="margin"
<code>\work{...}</code>	A work heading, title and size	div type="work" with head
<code>\sketch{...}</code>	Edward Hopper's ink record drawing stands here	figure type="record-sketch"
<code>\clipping{...}</code>	Something printed and pasted to the leaf	figure type="clipping"
<code>\ink{colour}{...}</code>	The colour a cell is written in, where it means something; classified from the photograph by a script, never by the language model (Section 7.2)	hi rend="ink-red" etc.
<code>\entry{as written}{ISO}</code>	A diary entry begins; the writer's date and the assigned date kept apart	div type="entry" with head/date@when
<code>\ruledoff</code>	The writer's rule under a sum, in the cash book	row rend="rule-above" on the row beneath
<code>\keywords{...}</code>	The batch's index terms, one facet each	keywords in textClass
<code>ledgertable</code>	The ruled columns	table, row, cell

```

<pb n="95" facs="https://resourcespace.whitney.org/pages/view.php?ref=17982" xml:id="sheet
-17982"/>
<p><seg hand="#jo">1942</seg></p>
<figure type="record-sketch"><head>Ink record sketch by Edward Hopper</head></figure>
<div type="work">
<head type="work"><« Night Hawks » 33 x 60</head>
<p><seg hand="#jo">Finished Jan 21, 1942</seg></p>
<p><seg hand="#jo">Blockx and Winsor & Newton colors, W & N. Zinc white, poppy oil.
English linen, domestic priming.</seg></p>
<p><seg hand="#jo">Night + brilliant interior of cheap restaurant. Bright items; cheery
wood counter + tops of surrounding stools; [...] Sign across top of <del>shop</del>
dark - Phillies 5¢ Cigar, picture of cigar.</seg></p>
<note place="margin">curving - above « window », with a caret beneath the line</note>
<note place="margin">restaurant - above the struck « shop »</note>
<note type="editorial" resp="#editor">this is the other side of the transaction batch 2
transcribed at leaf 25, where the Art Institute buys Compartment C in 1938 and
returns it in 1942 in part payment for Night Hawks. [...]</note>
</div>

```

5 The TEI encoding

5.1 Why the TEI is derived and not written

The TEI files are produced by a script from the transcriptions alone. The script re-reads no sheet, and the method requires that it must not: each macro maps onto a fixed TEI construct, the conversion is mechanical, and a difference between a transcription and its TEI is by construction a defect in the converter and not an editorial decision. This is the reverse of the usual division, in which the TEI is the source of record and a display is derived from it, and it was chosen for three reasons.

The first is the model. Asked to write TEI it does so fluently, and its fluency is the problem: well-formed markup, plausible structure, invented attribute values, and no structural signature to the failure. A sixteen-macro language has a small surface, each macro answers one editorial question, and a macro used wrongly is a visibly wrong answer rather than a subtly wrong element.

That decision has a precedent worth naming, because it makes the fourteen macros an editorial move rather than an idiosyncrasy. Cummings and Willcox (2013), digitising the Stationers' Register of 1557 to 1708 with a keying company that charged by the kilobyte, defined `tei_corset`: a deliberately reduced customisation of thirty-four elements with abbreviated names and constrained attribute values, not itself TEI-conformant, which XSLT converted up into conformant TEI while promoting presentational features into semantic markup. The shape is the same as this edition's and the motive is not: they shrank the markup to lower a bill, these macros to bound what a model may write; the consequence is shared, that an intermediate language must be converted by something testable, and the test is the schema.

The second is the checker. The transcription is read by people beside the photograph, in a rendering produced from the same file, and a LaTeX source with sixteen macros is also readable in itself, as a diplomatic transcription in a printed edition is, and can be corrected with a text editor.

The third is validation. Because the TEI is derived, it can be validated on every build against a customisation that declares exactly the elements the converter emits, and a validation failure

identifies a converter bug or a macro misuse without any ambiguity about whose judgement is at fault.

5.2 The customisation

The customisation is an ODD, `hopper.odd`, in the form the Guidelines define for one (TEI Consortium 2026), served at the URL every exported file names, from which both the schema and its documentation derive. It declares fifty-nine elements, and the fifty-nine are not a sample of TEI's six hundred but the exact reach of the conversion: every element exists because a macro maps onto it. The modules and elements included are shown in Table 5. Seven attribute value lists are declared closed, so that a `@type` on a division, a figure, a head, a note or a term, and a `@rend` on a highlighted cell or a ruled row, is drawn from a list and not invented per file, and the validator refuses any other value. The value lists are: division types (`leaf`, `part`, `work`, `entry`); figure types (`record-sketch`, `clipping`); head types (`work`); note types (`editorial`); the keyword facets (`medium`, `place`, `work`, `person`, `dealer`, `collection`, `society`, `publication`, `prize`, `client`, `subject`, `feature`); renditions on a cell (`italic`, `bold`, `mono`, `ink-red`, `ink-pencil`, `ink-blue`); and the one rendition a row may carry (`rule-above`).

Table 5: The modules and elements the ODD includes; nothing else validates.

Module	Elements included
textstructure	TEI, body, div, text
core	add, author, bibl, date, del, gap, head, hi, item, lb, list, name, note, p, pb, publisher, resp, respStmt, term, title, unclear
header	availability, catDesc, category, change, classDecl, edition, editionStmt, editorialDecl, encodingDesc, fileDesc, handNote, idno, keywords, profileDesc, publicationStmt, revisionDesc, sourceDesc, taxonomy, teiHeader, textClass, titleStmt
msdescription	history, institution, msDesc, msIdentifier, origDate, origin
figures	cell, figure, row, table
transcr	handNotes, supplied
linking	seg

The header of each file carries the provenance that the method requires. The two authors are named in the order the inside cover gives them; the responsibility statement names the model that read the sheets and the date, and states that the pass is a first machine pass unchecked by a person; the availability statement records that the transcribed text is in copyright and that the file is an unauthorised working document; the manuscript description carries the museum's accession number and the volume's dates; the hand notes declare the four hands; and the text classification carries the keywords, typed by facet.

```

<titleStmt>
  <title>Artist's ledger – Book II, batch 6</title>
  <author>Josephine Nivison Hopper</author>
  <author>Edward Hopper</author>
  <respStmt>
    <resp>Transcription (first machine pass, unchecked against the sheets by a person)</resp>
    <name>Opus 5 (claude-opus-5), 2026-09-06</name>
  </respStmt>
</titleStmt>
[...]
<handNotes>

```

```

<handNote xml:id="edward"><p>Edward Hopper</p></handNote>
<handNote xml:id="jo"><p>Josephine Nivison Hopper</p></handNote>
<handNote xml:id="later"><p>A later hand, after 1967</p></handNote>
<handNote xml:id="unidentified"><p>Unidentified</p></handNote>
</handNotes>

```

5.3 Tables, ink, and the cash book

The ruled tables are encoded as tables, one row per ruled line and one cell per column, with an empty label row where the leaf heads no column. The colour of ink, where the pass that reads it has run, is encoded as a rendition on the highlighted cell; a rule the writer drew across the money column above a row is a rendition on the row. The following rows from Book IV leaf 121 show a charge in black, a receipt in red, a sum in pencil and a ruled subtotal, and one receipt whose day and word could not be read.

```

<row><cell>June 12</cell><cell>Frank K. M. Rehn</cell><cell></cell><cell></cell></row>
<row><cell></cell><cell>Sattillo Roof Tops w.c</cell><cell>750</cell><cell>00</cell></row>
<row><cell></cell><cell>less commission</cell><cell>250</cell><cell>00</cell></row>
<row rend="rule-above"><cell><hi rend="ink-red"><unclear>12</unclear></hi></cell>
  <cell><hi rend="ink-red"><unclear>rec'd</unclear> <gap reason="illegible"/></hi></cell>
  <cell>500</cell><cell>00</cell></row>
<row><cell><hi rend="ink-red">12</hi></cell><cell><hi rend="ink-red">rec'd by check</hi></cell>
  <cell><hi rend="ink-red">500</hi></cell><cell><hi rend="ink-red">00</hi></cell></row>
[...]
<row><cell></cell><cell><hi rend="ink-pencil">1800 + 560</hi></cell>
  <cell><hi rend="ink-pencil">53</hi></cell><cell></cell></row>
<row rend="rule-above"><cell></cell><cell></cell>
  <cell><hi rend="ink-pencil">2360</hi></cell><cell><hi rend="ink-pencil">53</hi></cell></row>

```

The two red rows are one receipt: the writer wrote it across the row that carried the subtotal, rubbed it out, and wrote it again below, and the transcriber's note on the leaf says so. The accounts reader counts it once. This case, which recurs on nine leaves of the volume, is the clearest instance of why the transaction is recovered from the document by a rule rather than encoded in it: the document records the receipt twice, and the fact that it is one receipt is an inference.

5.4 The notebooks

A diary entry is a division of type entry whose head carries the writer's date as written and the assigned date in the @when attribute. The two are kept apart because the assignment is the one editorial claim an entry carries: *Garrulities* has entries whose day is a blot and whose year is inferred from the run of dates around them, and the encoding must not present an inference as a reading. The following is the black notebook's page 64, the one sentence of Edward Hopper's stated purpose that the ten books preserve, with the writer's own recasting of it below and the transcriber's note distinguishing the two.

```

<pb n="64" facs="https://resourcespace.whitney.org/pages/view.php?ref=98119" xml:id="sheet-98119"/>
<div type="entry">
  <head><date when="1946-01">Jan 1946</date></head>
  <p><seg hand="#jo">Jan 1946. <hi rend="italic">Quoting E. H.</hi> <lb/>
  Cruelly pinned down to definition <lb/>
  of his purpose generally, I extract: <lb/>
  " the arrestation of a moment <lb/>

```

```

in time - with the most acute <lb/>
realization.</seg></p>
<p><seg hand="#jo"> to arrest <lb/>
    a moment in time <lb/>
<del>with</del> that moment <lb/>
most acutely realized. "</seg></p>
<note type="editorial" resp="#editor">both blocks are enclosed in red crayon – the first
from « I extract » down, the second entire – and the second is her own recasting of the
first, not a second saying of his.</note>

```

5.5 Validation

The converter's output is validated in continuous integration against the schema derived from the ODD, and the workflow fails on any invalid file. The validator exists because of a defect it found: for as long as the export was checked only for well-formedness, every file in the corpus was well-formed XML and invalid TEI, since the converter had placed a head at each leaf inside a single division per batch, and TEI places a head only at the start of the division it heads. Nothing had said so because nothing had asked. The lesson is general and is recorded in the validator's own header: a deposit of unvalidated TEI is a deposit of XML.

6 Interface and reproducibility

6.1 The reading beside the photograph

The interface at <https://hopper.commutator.io> was built for this project. It renders the transcription in a left pane and the Whitney's photograph of the sheet in a right pane, and scrolling the transcription past a page-break element turns the photograph to that sheet, as facsimile viewers generally do; what is particular to it is that the photograph is the museum's own file, fetched as it is looked at and never stored, and the citation described below. Every doubtful figure can be checked against the hand it was read from without a second window. A Cite control on each sheet builds three citations from what the page already knows, a sheet, a leaf and a batch, each naming the museum's object clause and the edition's reading clause separately, and offers only those that exist: a sheet the Hoppers did not number gets no leaf citation, because inventing a number for the sake of a tidy citation would invent a fact about the object. The citation names the model and the date of the reading, because a citation of this edition cites a reading and not a fact, and a reading that could not be dated could not be superseded.

The reading views for the notebooks work the same way and are addressed by the museum's reference rather than by a page number, since an opening photographed whole carries two page numbers and a cover none.

The interface was built before the transcription was finished, and the order matters, because a reading that can be seen beside its witness is a reading that can be debugged. A misread cell in a file of ruled rows is invisible: it is a well-formed number in a well-formed column. It becomes visible in two places the site provides, the pane where it stands beside the photograph, and the drawings the derivations make of it, where a wrong figure is an outlier and a wrong year a hollow. Most of the corrections the repository records were found in the second place and settled in the first. The rule that a red figure in Book IV is money received, applied to the volume's opening leaves, turned eleven prices into receipts; the rise showed in the accounts before it showed in any file, and reading leaf 7 beside its photograph showed the

convention beginning on that leaf's last line. Three ruled rows under one title on Book I leaf 80 read as three sales until the leaf, in the pane, showed a sum written above them and the word «payment» in each; the accounts now book one sale of 2500 paid in three parts. The absence of 1921 from the cash book was a hollow in the calendar before it was a year line misread on leaf 53. None of this needs an elaborate interface, only the reading and the witness on one screen and the derived figures on the next, from the first batch onward; an edition that builds its reading view last finds its errors last.

6.2 Derived data, and the rule that a build must fail

Everything the interface shows beyond the transcription is derived from the transcription by a script whose header states what it counts and what it therefore is not, and the continuous integration workflow regenerates each derived file and fails if the committed copy differs. The manifest of what is transcribed is observed from the files, never declared, so that a progress figure cannot claim work that was not done; the only declared states, *checked* and *skipped*, are held in a separate file that is empty. The citation metadata is generated from a census of the files, on the ground that a hand-written sheet count is wrong from the next batch onward. The glossary of abbreviations and notations is hand-declared, but each entry carries a pattern that must match a transcribed leaf, and the search runs over the writer's words and never over a transcriber's note, so that an entry cannot cite as evidence the sentence written because of it.

The accounts reader recovers transactions from the ruled rows. It counts a row only where the leaf puts a price in one cell and a date in another, since the column a date stands in says whether it dates the sale or the cheque; on Book I's etching leaves the first column dates the exhibition, so those rows are set aside rather than dated wrongly. Where a row states a price, a rate and a receipt, the receipt is checked against the price less the commission; of ninety-five such rows, ninety agree to the cent, and the five that do not are reported and not corrected. For Book IV the reader has a second set of rules, set out period by period in a method note, because the writer's form changes four times across the volume: items are the charge; a bill counts only where nothing was itemised since the last settlement; a bare figure is a subtotal where the arithmetic of the lines above it says so, since the leaf rules a line above a subtotal and the transcription does not record the rule; a «less» line is a deduction whichever thing it names; a title charged more than once with a payment on account is one picture at the largest figure. The check on the reader is the book's own pencil totals, and the reader's receipts reconcile to them to the cent wherever the leaf's arithmetic closes.

Two of the reader's rules were found by bugs, and the method note records them: an unanchored match on the word *bill* took a story title, «Wild Bill in Deadwood Gulch», for an invoice and lost a drawing of 1915; a header rule read «May 25th Initial Letters» as a new client and counted a bill of 225 twice. Both were caught by comparing per-year totals with leaves read by hand. A faithful transcription can surface a reader bug without containing one.

6.3 The rules, stated as proofs

A reader whose rules are written in prose and applied by a script can drift from them unnoticed, and a total can reconcile by accident. The derived accounts are therefore also restated in the Lean 4 theorem prover, in a directory of the repository that depends on nothing but Lean, so that the continuous integration build fails if the script and its stated rules disagree. The formalisation was written by Claude Opus 5 under the author's direction, on 14 September

2026. Every figure in the ledgers is dollars and cents and every commission a third, a quarter, a tenth or fifteen per cent, so money is represented as integers in sixtieths of a cent, where all of them are exact. General lemmas are proved for every input by decision procedures for linear arithmetic; the data is checked by evaluation in the kernel (`decide +kernel`) and never by compiled code; and warnings are errors, so an unfinished proof fails the build.

Two certificates are proved. The first concerns the accounts as published: each activity's years are recomputed from its rows, and the identity that what accrued equals what was collected plus what is outstanding plus what was not recorded is derived from the rows rather than read off the totals; the reader's two-cent tolerance is proved to accept every figure written to the cent after a truncated third and to reject a slip of a third of a dollar; the ninety-five checkable rows are proved to yield ninety agreements and exactly the five disagreements the reader reports; and each hand-declared pencil total is proved equal to the reader's receipts for that year plus the declared omissions. The second concerns the Book IV reader. Its four passes are stated as Lean functions, and the properties the method note asserts in prose are proved of them for every input: a bill is charged only where nothing was itemised since the last settlement, so no commission is counted twice; a «less» line is never a charge; a bare figure is a subtotal exactly when the answering line or the arithmetic says so; a picture paid in instalments is charged once; the year does not carry across a gap between batches. The kernel then evaluates the functions over all 3,281 transcribed rows of the volume and proves that each row is decided exactly once and that the Lean reader and the JavaScript reader agree on every row's kind, year and disposition: the JavaScript exports the lexical features it recognised and the decisions it took, Lean recomputes the decisions from the features, and altering a single disposition in the export fails the build.

The exercise refined one rule — a «less» line that dittoes the receipt above it reads as that receipt — and turned the missing year of Section 6.1 into a theorem written to fail once leaf 53 is corrected at source. The limits are the method's own: the certificates say nothing about a reading, a figure misread consistently in the rows and in the sums would pass, the regular expressions that recognise a row's words are inputs rather than proved, and what is proved is that the script agrees with its stated rules and not that the rules read the writer's practice rightly. They are checked on every push, and an equivalent GitLab CI configuration is provided.

6.4 What the data shows, as an illustration of the condition

Figures 2 to 5 reproduce four of the interface's drawings from the derived files, and are given here not for their findings but to show what a fully derived figure looks like: each names its source file and, where that file records it, the date it was generated, and each is regenerable by anyone with a clone of the repository.

The format census reads only the size in the title line, in Edward Hopper's hand: 199 of 270 work headings state one, they make 78 formats, 72 of them wider than tall and none square. The materials census reads only the paint formula in Josephine Hopper's hand, counted once per work and never from a transcriber's note: 77 headings carry one, and the succession they state (the caption of Figure 3 gives it) is dated by the leaves themselves for 43 works and by nothing else. The money is given three ways, as the caption of Figure 4 enumerates, because the work books and the cash book are different records and should not be added. The days on which the books record a transaction number 1,203 over fifty-four years, Sundays carrying a sixth of a weekday's count, with a hollow every summer, which is the Cape; the drawing is




Figure 2: Every format, plotted: width across, height up, one dot to a format and its area standing for how many works were made at it, 78 formats from 199 work headings that state a size. The diagonal is where a picture would be square; 190 of the 199 works lie below it, 9 above, none on it. The etching plates are in the corner, the watercolour sheets through the middle, the late oils at the top right. After the Technique tab of hopper.commutator.io. Drawn by Claude Fable 5.1, from `formats.json` (undated).




Figure 3: A succession, not a preference: each material against the years the leaves date it to, earliest first, one dot to a year and its size standing for how many works that year names it. Rembrandt colours and zinc white in poppy oil at the start; Winsor & Newton and flake white in linseed oil and turpentine from the middle 1940s; the canvas turning from National to Herga and the priming from double to single across the 1950s. Only the 43 works dated by the leaf itself are plotted. After the Technique tab. Drawn by Claude Fable 5.1, from `materials.json` (undated).




Figure 4: The money, three ways: what the pocket book receipts each year from 1920 (dark bars), what the work books state as the net of sales dated in that year (line), and the illustration receipts of the first decade (pale bars). The two art series are not the same number and should not be: one is the dealer’s cheques as they arrived, the other the studio’s record of what left. Every figure is a floor, read from the transcribed rows under the rules of the accounts chapter. After the Accounts tab of `hopper.commutator.io`. Drawn by Claude Fable 5.1, from `accounts.json` (2026-09-19).

captioned to say that a dense month is a month of transactions and not of painting, because a grid of squares reads as productivity and here that is backwards.

7 Discussion

7.1 The reading against the image

The transcription is a first pass and has not been checked by a person. For this article, the two quotations on which the edition’s own reading of the archive most depends — the Goethe line closing Book III leaf 19, and Edward Hopper’s one statement of his purpose in the black notebook — were compared with the museum’s photographs by the model itself on 13 September 2026, at the published resolution of 1292 by 2000 pixels; Table 6 gives the result cell by cell. Both readings hold in every word, and the comparison found one omission: a small mark before the struck word on the notebook page, which the transcription does not record. The comparison is reported for what it is. It was made by the same kind of reader that made the transcription, it covers two of the 667 sheets transcribed by 13 September 2026, and it establishes only that a second reading agreed with the first. The edition reserves the word *checked* for a comparison made by a person, and the comparison is recorded in the transcripts as a note that says who made it and that it is not a human check.

A more systematic comparison exists for one volume and shows the limit of the reading. A pass that aligns each transcribed row of Book IV with the written lines of its photograph, in order to read the ink colour the model cannot see, reports seventy-four leaves, about half the volume, with at least one written line the transcription does not have. Many are the alignment’s own noise, and each must be looked at; at least one, an entry of 26 February 1917 on leaf 21, was a real omission. The consequence is bounded and can be stated: a missing line changes a sum, so the cash-book figures are floors in a second sense, and no such pass has run



Figure 5: The days the six books record, one row to a year from 1913 to 1967 and one cell to a month, the darker the more days that month carry an entry: 1,203 days in all, from 2,143 dated rows. A dense cell is a month of transactions, not of painting; the hollow every summer is the Cape, with the dealers behind them in New York. By day of the week: Sun 34, Mon 197, Tue 197, Wed 179, Thu 214, Fri 207, Sat 175, Sunday a sixth of a weekday. After the calendar of the Timeline tab, which draws every day; this draws every month. Drawn by Claude Fable 5.1, from `streak.json` (2026-09-19).

Table 6: Transcription against image, for the two sheets compared on 13 September 2026. Sheet refs 98119 (black notebook, page 64) and 16838 (Book III, leaf 19).

Sheet, place	As transcribed	As seen in the photograph
98119, head	«Jan 1946. <i>Quoting E. H.</i> »	Agrees; the second phrase underlined
98119, l. 2–3	«Cruelly pinned down to definition / of his purpose generally, I extract: »	Agrees
98119, l. 4–6	«" the arrestation of a moment / in time - with the most acute / realization. »	Agrees; «acute» underlined; block enclosed in red crayon from «I extract» down
98119, block 2	«to arrest / a moment in time / with that moment / most acutely realized. " »	Agrees; «acutely» underlined; a small cross before the struck «with» is not recorded
16838, title	«October on Cape Cod. 26 x 42. Finished on Cape Cod - in studio - Oct. 1946. »	Agrees
16838, margin	«To Rehn Gallery Nov. 26, 46 » at the left	Agrees
16838, paint	«Windsor + Newton linen canvas + Colors; flake white, linseed + turpentine. »	Agrees
16838, interlined	«bearberry» above «herbage»; «low» above «clumps»; «pale» above «blue at top»	Agrees, all three
16838, close	«"Peace, quiet, no birds sing." "Some day you will be quiet too" - E. H.'s favorite Goethe. (strictly off the record) »	Agrees
16838, sale	«3000 - 1/3 = 2000 June 14, 1956. »	Agrees

on the other volumes or the notebooks, about whose omissions nothing is therefore known.

7.2 Limits of the model as a transcriber and as an encoder

The experiment identified five limits, none of them a matter of error rate.

The model does not see colour reliably. Book IV's meaning is in its ink, and fourteen batches were transcribed without recording it; the colour was recovered afterwards by an image-processing script that classifies each money cell and writes the result into the transcription, on the ground that colour is the one thing on the leaf a machine can see better than the hand it is written in. That recovery is partial and its extent should be stated: on 20 September 2026 thirty-two of Book IV's one hundred and sixty-one sheets carry an ink mark, the pass having been run where a leaf's rows and the written lines of its photograph align exactly, and stopped where they do not. The stop is mechanical rather than palaeographic, so the remaining leaves await the human reading rather than a better classifier; where the convention itself begins, and how that was found, is told in Section 6.1. A method that depended on the language model for the ink would have produced a cash book in which a charge and a receipt cannot be told apart.

The model carries context forward as a rule, and rules fail at boundaries. The year of a Book IV entry is stated once, on a line of its own, and carried down the leaf; the year line of leaf 53 (Section 6.1) is the case, where the transcription keeps the line as written and the reader carries the wrong year. That is the intended behaviour of both, and it is wrong, and it is listed for correction at source rather than repaired in the derived data.

The model stretches a macro to the nearest meaning, and the apparatus has to be read for

where it did. \add was declared as the transcriber's supply and mapped to supplied; read in context, a majority of its 329 uses were Josephine Hopper's own insertions — a buyer's name over a struck one, a «later sent \$5» under a line — which TEI encodes as add in the writer's hand. Given one macro for an insertion, the model used it for both kinds, and the converter credited her afterthoughts to the editor.

The two senses are now two macros, and the corpus was divided between them by a rule checkable in the file rather than by a fresh reading: an insertion glued to the word it completes, or standing in a table's header row where the writer ruled a column and never labelled it, is the transcriber's; one standing free in a cell is hers. That assigns 137 uses to \supplied and leaves 192 as \add. The residue is honest rather than certain: a free-standing supply still reads as the writer's, and only the photographs will settle those.

The model reads names as language. The archive spells Crowninshield three ways — Crowninshield, Crownenshield, Crowneshield — and Saltillo two, and the model, left to itself, would settle them. The rule that nothing is normalised holds it to the leaf; the discipline is the rule's and not the model's.

The model is fluent in TEI, and that fluency is the hazard Section 5.1 describes. The decision not to let the model write TEI at all, and to derive the TEI mechanically from a seventeen-macro language it can hardly misuse, is the central encoding decision of the edition. Its cost is that the encoding is only as expressive as the macro set, and the macro set was designed for what these ten books contain; the transactionography of Tomasek and Bauman (2013), for instance, has no place in the transcription and lives entirely in the reader.

The model is confident about hands. The attribution of a hand on the letterforms is the one judgement in the apparatus that no arithmetic can check, and the model attributes readily. That difficulty is not the model's alone: Stokes (2015) argues that computational measures of similarity cannot on their own establish that two specimens were written by one person, since there is rarely any independent ground truth to verify the claim against, and that such output should direct an expert's attention rather than stand as a verdict. The rule that *unidentified* is a permitted answer, and the decision to attribute the whole of Book IV — a hand that predates the marriage by eleven years — to nobody, are the method's response. Where the edition does attribute, it does so on a document and not on the letterforms: the 205 record sketches of Books I, II and III carry @hand because the inside front cover of Book I declares, in Josephine Hopper's writing, that the drawings in the three books are Edward Hopper's, and the twenty-four sketches of Book V carry none, because that sentence does not name Book V. A human palaeographer with the physical volumes would do better than either, and the edition says so on the leaf.

7.3 Human validation

The edition distinguishes three states of a batch, and only one of them is observed from the files. *Drafted* is what a file's existence proves. *Skipped*, for a batch judged to hold nothing to transcribe, and *checked*, for a batch that a person has compared with the photographs sheet by sheet, are declared in a status file, and that file declares nothing for any batch. The reason is stated in its comment: a hand-kept status table drifts optimistic and nobody notices, so the two states no file can prove are the only two a person may write, and the site shows a transcribed batch as drafted because that is observed rather than declared.

What a human check consists of has been defined as an open question of the edition rather than assumed. The protocol under discussion is the one this article's comparison table per-

forms in miniature: a person opens the reading beside the photograph, reads each transcribed line against the sheet, corrects the transcription at source with a note that names who checked it, and only then may the batch be declared checked; the citation of every sheet in that batch then changes to say that a person has gone through it, and only a person may declare that. Until then, every citation of the edition names a machine pass and a date, and the address of the sheet does not move when the reading does.

7.4 The pace, and where the cost went

An account of this method that did not state its speed would be evasive, since speed is the whole of what changed. The figures are in the repository's own history and can be recovered from it.

The transcription of 667 sheets took nine days, from 5 to 13 September 2026, in sixty-two passes of up to twelve sheets. Nobody keyed a character: a pass is a photograph read and a file written, and the twelve sheets of a batch take the model minutes. In the same fortnight the repository received the scripts that derive the accounts, the formats, the materials and the calendar (some 11,000 lines), the site that places the reading beside the photograph (some 12,500 lines), the TEI converter and its validator, and later the 1,800 lines of Lean that certify the arithmetic. A transcription project of this size, keyed and proofed by people, is measured in a person's months; this was a fortnight of one person's attention and a model's time.

That comparison is the easy half, and it is not the interesting one. What the method actually changes is the second operation, the one that has no name in the traditional workflow: *restructuring*. Once the transcription is a constrained language rather than a document, every question one might put to the archive becomes a script of a few hundred lines — what was paid and by whom, which formats recur, when the colourman changed, on what days the books were written in. Each of those is a census this edition publishes, each was written in an afternoon, and each is regenerated on every build. The cost of asking a new question of the corpus has fallen by more than the cost of reading it, because reading is done once and asking is done for ever.

The cost did not disappear; it moved, and it moved to the two places this article has been describing. The first is the definition of rules. The Book IV reader changed its rule four times across the volume's fifty-four years, and every one of those changes is an editorial decision about what the writer meant that no model and no script could take: whether a bare figure is a subtotal, whether a «less» line is a deduction, whether three payments are three sales or one. Writing them down — in the method note, then as Lean functions proved for every input — took longer than transcribing the volume. The second is verification. 667 sheets are read and none is checked. A pass costs minutes and a human check costs an hour a batch, so the reading has produced, at the pace described above, a debt of some fifty hours that the method cannot discharge and does not pretend to.

The same arithmetic applies to the three sibling editions of Section 8, and states the size of what is left. The Grothendieck fonds stands at 1,937 pages in 122 passes, sixteen pages to a pass, against the sixteen thousand pages the edition's own catalogue counts as openly accessible at Montpellier: at that rate the reading is some nine hundred further passes, which is months of machine time and, at an hour a batch, about nine hundred hours of human checking — that is, a reading a machine can finish this year and a verification no individual can. The mathematicians' manuscripts in Gallica stand at 647 views in 35 passes, eighteen views to a pass, against 25,258 views catalogued: some 1,400 further passes. Neither number

is an argument for doing it; both are the argument of this article, which is that the transcription is no longer the bottleneck and the checking is, and that an edition which cannot say which of its sheets a person has read is publishing a machine's confidence as though it were scholarship.

There is a third cost the traditional account treats as fixed, and it is the apparatus itself. Everything in Figure 1 below the reading — the converters, the censuses, the validators, the site, the workflows that rebuild and compare — is code, and code of that kind was the part of an edition that took a developer and a grant. Here it was written by the same model that read the sheets, through a code assistant, under one person's direction, at a pace the repositories' histories give. This edition's apparatus, some twenty-three thousand lines of scripts and interface and eighteen hundred of Lean, was written between 5 and 20 September 2026 in some two hundred and twenty commits. The siblings show what an edition costs once the schema exists: the Grothendieck edition, first and therefore longest, 14,500 lines in 246 commits over six weeks; the Gallica edition, 11,800 lines in 56 commits over ten days; the morphology course, 5,700 lines in 104 commits over two days, the schema by then known and only its values new. Lines and commits measure activity, not quality; what they establish is that a stated schema can be filled in for a new corpus in days.

That ease has two consequences, both met here. Code written at that pace is written faster than it is read, and what makes it trustworthy is the discipline described for the transcription: every derived file regenerated and compared on every build, every rule stated where a reader can find it. Two of the reader's rules were found by bugs (Section 6.2), and the calendar of Figure 5 went stale because its generator wrote to a path no build checked. The assistant does not remove the checks; it makes them the whole of the guarantee. And when a converter costs an afternoon, what remains expensive are the editorial decisions — what a bare figure is, whether three payments are one sale — which the code cannot take: the cost of an edition moves toward its rules and its checking, which is where it belongs.

7.5 What the method permits, and what it does not

The method permits a reading of a seven-hundred-sheet archive to be produced in days, cited by leaf and reference, exported as valid TEI, exposed to correction beside its source, its derived data regenerated under published rules and its rules proved of the script that applies them.

It does not produce a checked transcription, does not see ink, does not attribute hands with authority, does not encode transactions, and cannot produce a text more expressive than its macro language. And it does not remove the need for the person it was designed around, whose correction is the only event that changes a batch's state.

8 The method against four corpora

The apparatus described here was not built for this corpus alone. Three sibling editions run on the same tooling — the same restricted LaTeX, the same rule that an unrecognised construct fails the build, the same reading placed beside its witness — and they were chosen to differ in the two things that matter: what the document is, and what kind of witness stands behind it. Two are archives of art, two of mathematics; three are written and one is spoken. Table 7 gives their state, read from the four repositories on 19 September 2026.

Figure 6 sets out what may be asked of a machine in each of them, and the answer is not a matter of the model's quality: what moves the line between the model, the scripts and the person is the document.

Table 7: The four corpora the method has been run against, with the address at which a reading can be checked.

Corpus	Holder	Witness	Verifiable address	Transcribed
Hopper ledgers and notebooks	Whitney Museum	Ruled tables, two hands	ResourceSpace reference	667 sheets of 673
Grothendieck fonds	University of Montpellier	Mathematical prose and notation	Folder and page	1,937 pages in 63 folders
Mathematicians in Gallica	Bibliothèque nationale de France	Manuscripts, sixteenth to nineteenth centuries	Gallica view (IIIF)	647 views in 10 volumes
Jean-François Debord's morphology course	Beaux-Arts de Paris, published by PSL	A filmed lecture, with automatic captions	Timestamp in the session	45 sessions, 10,810 timestamps




Figure 6: What may be asked of a machine, by what the source is. The division of labour is the same in all four editions — the model reads, the scripts derive, the schema validates, the person checks — and what moves between them is the document: a ruled table is read cell by cell and its colour not at all; a filmed lecture is transcribed by the platform before the model sees it, and what the model adds is an index addressable to the second. The editorial condition does not move with them: nothing is asserted that a reader cannot address.

Table 8: The schema of Figure 1, with each edition’s values, read from the four repositories on 20 September 2026.

	Hopper	Grothendieck	Gallica	Debord
Witness	Photograph, fetched at view time	Facsimile, streamed from Montpellier	IIIF view, mirrored and git-ignored	Recording and its captions, downloaded, never versioned
Model’s task	Reads the cells	Reads, then restates the mathematics	Reads the page	Indexes a reading the platform made
Reading	batch-NN.tex, twelve sheets a pass	batch-NN.tex and .modern.tex, up to twenty pages a pass	batch-NN.fr.tex and a modernised file, up to twenty views a pass	A lexicon of terms, inflections and the recogniser’s mistakes; references and digressions
Derivations	TEI, PDF, HTML; accounts, formats, materials, days	TEI, PDF, HTML; folder tags from the keyword line	TEI, PDF, HTML; the typed catalogue	Occurrences: identifiers and seconds, no word of the course
Verifications	Schema; Lean; regenerate and compare	Schema; the two registers kept in step; the linked editions still served	Schema; the sources still served	The index’s tests; two aligned texts, accents kept
The person	Empty	Declared in a status file	Declared in a status file	Listens at the second; the verdict enters the lexicon

Table 8 makes the same point station by station: the five stations of Figure 1 are filled in four times, no edition adds one or lacks one, and the values are set by the document — a photograph, a streamed facsimile, a mirrored view, or a recording nobody may republish; a model that reads, that reads and restates, or that indexes what another machine read; a person who compares against a photograph in three editions and against a sound in the fourth.

What the three siblings taught the method is specific in each case. The Grothendieck fonds carries a second register the ledgers refuse: its transcriptions come in pairs, a diplomatic file and a modernised one (124 and 53 files), because mathematical notation of 1960 is not read today as it was written, whereas «30 - 1/3» means in 2026 what it meant in 1927 — the decision against a second edition here (Section 4.3) is a decision this project had to take twice, in opposite directions. The Gallica corpus made the address itself the unit: a volume enters only where the library serves the whole of it, and a manuscript that Gallica merely references on a partner's site is in no table and no thumbnail, which is the same rule this edition states about documents with no public digitisation. And the morphology course made the witness something other than an image.

That last case is the one that generalises the claim. Jean-François Debord's course is seventy-seven hours of speech; the first transcription is not the model's but the platform's automatic captions, and the model's work is an index — 95 terms, 73 references and 139 digressions, each carrying the second at which the thing is said. The repository publishes no caption and no video: only video identifiers, term identifiers and offsets in seconds, the captions themselves staying out of version control, as the photographs stay out of this one. The editorial condition is unchanged and its instrument moves: a reader does not compare the reading with a photograph but with the recording, at the timestamp given, by listening. What every corpus here has in common is not a document type; it is that no sentence is asserted which a reader cannot address.

The platform corpus also shows where the condition begins to fail, and it is worth stating before such corpora multiply. A sheet at the Whitney will answer to the same reference in twenty years; a video can be withdrawn, re-edited, or have its captions regenerated by a newer recogniser that yields different words at different seconds. The address survives and what it serves does not, and nothing in the citation announces the change. Three things follow, none of them solved here, and the conclusion returns to them: such a citation must carry the date it was consulted and a fingerprint of what was consulted — a hash of the caption file, a duration, a count of segments — to prove that the witness moved without republishing it; a verification of this kind is revocable, and an interface that keeps a dead timestamp on screen asserts what it can no longer support; and a platform forbids the mirror while promising no permanence. An edition may not copy what nobody undertakes to keep.

9 Conclusion

This article has described an edition in which a language model is the first transcriber and the editorial apparatus is designed to make its reading checkable rather than trustworthy. The method rests on four decisions: a restricted transcription language whose macros answer one editorial question each; a mechanical derivation of TEI from that language, validated on every build against a customisation that declares exactly what the converter emits; a derivation of all counts and figures from the transcription under published rules, checked against the writer's own sums and, for the accounts, certified in Lean against the rules as stated; and an interface

that places the reading beside the photograph and cites every sheet by the museum's own reference and the edition's own dated pass.

What the experiment opens is a division of labour. The model reads; the scripts derive; the schema validates; and the person checks, corrects at source, and declares. The first three have been done for this corpus; the fourth has not, and the edition's status file, which declares nothing, measures where the method stops. The corpus is deposited under a persistent identifier. The human check that turns a drafted batch into a checked one is to be settled with the scholars who will do the reading, is tracked as an issue of the repository, and is the one on which every claim in this article depends.

What it also opens is the range of what can now be read. The four corpora of Section 8 are all archives in the ordinary sense — ledgers, manuscripts, working papers, a filmed course — and between them they already cross the line that used to bound this kind of work: three are written and one is spoken, three are held by institutions that publish facsimiles and one is held by a platform that publishes none. Behind that line lies the larger part of what the twentieth and twenty-first centuries have left: recorded lectures and interviews, broadcast archives, the video collections of national libraries, and the material of social platforms — posts, threads, the captions that accompany millions of hours of speech — most of it never read in the sense this article means, because reading it by hand was never conceivable.

The method described here does not become easier on that material; it becomes more necessary, and its conditions become harder to meet, in the three ways Section 8 sets out: the address may outlive what it serves, the rights forbid the mirror while promising no permanence, and the first transcription is often the platform's rather than the editor's, which turns the model from the first reader into the second and makes the provenance statement more important rather than less.

None of that changes the condition this article has argued for. An edition is worth what its reader can check, whether the checking is done by looking at a photograph, by adding up a column, or by listening at a timestamp. What changes is how much there is to read, and how badly it needs somebody to say, of each piece of it, who read it and on what day.

About the author

Michel Hua is an independent consultant at Commutator, Paris, working with large organisations on data and cloud applications: platform migrations, data governance, and the adoption of generative AI. The working assumption there is that failures occur at the interfaces between systems rather than inside them, and that a deliverable must be legible enough for the client's own engineers to take over. The editions described here apply that discipline to archives institutions have digitised and nobody has read: the Hopper ledgers, the Grothendieck fonds, the mathematicians' manuscripts in Gallica, and an index of a filmed course. michel@commutator.io.

Declaration of Generative AI and AI-assisted technologies in the writing process

During the preparation of this work the author used Claude Opus 5 and Claude Fable 5.1 (Anthropic), through Claude Code, in order to transcribe the digitised sheets of the corpus into the restricted transcription language described in Section 4, to compare two sheets with their

photographs as reported in Section 7.1, to write the scripts, the interface and the continuous-integration workflows described in Sections 5 and 6, to write the Lean formalisation of the accounts described in Section 6.3, and to draft this article from the transcriptions and the derived data under the author’s direction. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the publication. The transcriptions have not been checked against the photographs by a person, and the article says so wherever a claim depends on them.

Data availability

The transcriptions, the TEI export, the ODD, the scripts, the derived data, the Lean proofs, the interface and the sources of this article are in the public repository at <https://github.com/Commutator-I0/hopper-project>; the interface is served at <https://hopper.commutator.io>. The software and the derived data are released under CC0 1.0; this article under CC BY 4.0. The transcribed words of Josephine and Edward Hopper are © Heirs of Josephine N. Hopper, licensed by Artists Rights Society (ARS), New York, and are quoted for criticism and study; the photographs are the Whitney Museum’s and are not stored by the edition. The corpus is deposited with a persistent identifier: the concept DOI <https://doi.org/10.5281/zenodo.22770550>, which resolves to the most recent archived version, with a Software Heritage snapshot of the repository’s history beside it.

References

- Crosilla, Giorgia, Lukas Klic, and Giovanni Colavizza. 2025. “Benchmarking Large Language Models for Handwritten Text Recognition.” *Journal of Documentation* 81 (7). Preprint: arXiv:2503.15195.
- Cummings, James, and Pip Willcox. 2013. “Stationers’ Register Online: A Case Study of a Byte-Reduced TEI Schema for Digitization (tei_corset).” *Journal of the Text Encoding Initiative* 6. <https://doi.org/10.4000/jtei.926>.
- Humphries, Mark, Lianne C. Leddy, Quinn Downton, Meredith Legace, John McConnell, Isabella Murray, and Elizabeth Spence. 2025. “Unlocking the Archives: Using Large Language Models to Transcribe Handwritten Historical Documents.” *Historical Methods: A Journal of Quantitative and Interdisciplinary History* 58 (3). <https://doi.org/10.1080/01615440.2025.2500309>.
- Larson, Madeleine. n.d. *Josephine N. Hopper’s Diaries, 1933–1956: Typed Transcript*. 5 parts. Provincetown Art Association and Museum. Scans consulted through the museum, cited by part and typescript page.
- Levin, Gail. 1995a. *Edward Hopper: A Catalogue Raisonné*. New York: Whitney Museum of American Art and W. W. Norton.
- Levin, Gail. 1995b. *Edward Hopper: An Intimate Biography*. New York: Knopf.
- Lyons, Deborah. 1997. *Edward Hopper: A Journal of His Work*. New York: Whitney Museum of American Art and W. W. Norton.

- Scholger, Martina. 2020. "Taking Note: Challenges of Dealing with Graphical Content in TEI." *Journal of the Text Encoding Initiative* 12. <https://doi.org/10.4000/jtei.2800>.
- Stokes, Peter A. 2015. "Digital Approaches to Paleography and Book History: Some Challenges, Present and Future." *Frontiers in Digital Humanities* 2: 5. <https://doi.org/10.3389/fdigh.2015.00005>.
- Strutz, Sabrina. 2026. "A Multi-Dimensional Evaluation Framework for Assessing LLM Performance in TEI Encoding." *Journal of Open Humanities Data* 12 (1). <https://doi.org/10.5334/johd.484>.
- TEI Consortium. 2026. *TEI P5: Guidelines for Electronic Text Encoding and Interchange*. <https://tei-c.org/guidelines/p5/>.
- Tomasek, Kathryn, and Syd Bauman. 2013. "Encoding Financial Records for Historical Research." *Journal of the Text Encoding Initiative* 6. <https://doi.org/10.4000/jtei.895>.
- Whitney Museum of American Art. n.d. *Edward and Josephine Hopper Resources*. <https://whitney.org/research/edward-and-josephine-hopper-resources>.

A A complete transcribed leaf

The transcription of Book III leaf 19, October on Cape Cod (ref 16838), as the model wrote it in the restricted language, followed by the note recording the comparison with the photograph. The note is given verbatim, as the file carries it: the comparison of 13 September 2026 was made for a critical note on the ledgers that was abandoned, and the sentence saying so stays in the transcription because it is what the file recorded on the day, and the comparison itself is unaffected.

```
\sheet{16838}{19}

\sketch{}

\work{October on Cape Cod. \quad 26 x 42. \quad Finished on Cape Cod - in
studio - Oct. 1946.}

\marginal{To Rehn Gallery Nov. 26, 46 --- in the margin at the left of the title
and paint lines}

\hand{jo}{Windsor + Newton linen canvas + Colors; flake white, linseed +
turpentine.}

\hand{jo}{Beyond slatey grey road a strip of very pale (yellowish) sand across\
length of canvas. Low herbage digging in green. Tall grass beyond slightly\
pinkish with clumps of green. House fresh painted white, barn left dull\
Rooves of both slatey grey. Glimpse of red foundation, chimney dull red.\
Edge of locust grove L. dull green. Woods in background with green underfoot\
at front Edge. Sky wan, whitish with long strips of grey,. Some blue at top.\
"Peace, quiet, no birds sing." "Some day you will be quiet too" - E. H.'s
favorite Goethe.\
\quad (strictly off the record)}

\marginal{bearberry --- above « herbage »; low --- above « clumps »; pale ---
above « blue at top »}
```

\hand{jo}{Gone to Mr. + Mrs. O. W. Carpenter. Reprod. in Vogue, June 1954.
When sold?\\
R. H. Bresler Co. – payment caught up in Rehn Estate –\\
\quad Paid by John Clancy, new head of Rehn Gallery. 3000 – 1/3 = 2000
June 14, 1956.\\
Carpenter of Milwaukee – Heavy chain, bolt Co.}

\note{the marginal at the head is at the left here where leaves 13 and 15 put the same note at the right. « When sold? » is in a darker pencil than the line it follows and was added to it. The two quotations closing the description are hers to place and she places them: the first she gives no source, the second she gives to Goethe by way of Edward, and then she puts the whole attribution off the record. « favorite » is as written.}

\note{compared with the photograph at the Whitney's published size by Claude Fable 5.1 on 2026-09-13, for the critical note: the title line, the paint line, the description with its three interlineations, the two closing quotations, « E. H.'s favorite Goethe » and « (strictly off the record) », and the sale lines all read as transcribed. A machine's comparison, not the human check the edition reserves « checked » for}

B The derived tables

The tables below are regenerated from the edition's derived files on each build; those given here were generated on 13 September 2026, when 46 of the 48 transcription files existed. Every figure is a floor under a stated rule.

Table 9: Book IV, the illustration decade: what was charged and what was received, by year, as the transcribed leaves state it. Accrual at the leaf's own date column; a negative balance is a cheque with no charge line behind it. There is no row for 1921 because leaf 53's year line reads 1920 above entries of 1921 and the reader keeps the year as written; that year's income is counted under 1920.

Year	Charged	Received	Balance
1913	160.00	160.00	0.00
1914	1,225.00	1,225.00	0.00
1915	500.00	500.00	0.00
1916	991.00	991.00	0.00
1917	1,509.00	1,514.00	-5.00
1918	1,556.00	1,916.00	-365.00
1919	977.00	977.00	-365.00
1920	2,950.00	2,794.00	-209.00
1922	1,075.00	1,075.00	-209.00
1923	1,279.20	1,304.20	-234.00
1924	1,015.00	980.00	-199.00
1925	1,218.67	1,218.67	-199.00
Total	14,455.87	14,654.87	-199.00

Table 10: The work books (Books I, II, III and Dealers/Etchings): sales dated in the leaf's own date column, by year. Gross is the price written; the commission is the dealer's share as the row states it; net is what was left. Only rows that put a price in one cell and a date in another are counted, so the etchings are largely absent.

Year	Sales	Gross	Commission	Net	Receipted
1923	3	970	323	647	0
1924	6	4,500	1,500	3,000	0
1925	4	2,100	510	1,590	340
1926	11	5,275	1,758	3,517	850
1927	17	3,955	1,318	2,637	1,333
1928	24	5,774	1,919	3,855	1,197
1929	52	7,040	2,342	4,698	4,197
1930	34	18,378	6,126	12,252	3,000
1931	17	2,648	883	1,765	5,600
1932	4	3,880	1,293	2,587	2,400
1933	3	85	28	57	20
1934	13	3,680	1,227	2,453	800
1935	1	750	250	500	500
1936	6	4,420	1,473	2,947	267
1937	10	4,670	1,557	3,113	20
1938	11	4,990	1,648	3,342	0
1939	1	2,500	833	1,667	0
1940	14	6,182	2,061	4,121	667
1941	4	1,220	407	813	267
1942	3	2,780	927	1,853	0
1943	4	2,760	920	1,840	1,200
1944	5	4,850	1,617	3,233	0
1945	5	3,080	1,027	2,053	0
1946	6	7,650	2,550	5,100	2,267
1947	7	6,130	2,043	4,087	1,467
1948	4	7,500	2,500	5,000	1,500
1949	5	9,115	3,038	6,077	0
1950	12	4,130	1,377	2,753	1,053
1951	11	7,830	2,610	5,220	333
1952	3	7,750	2,583	5,167	2,667
1953	9	4,500	1,500	3,000	500
1954	6	14,780	4,927	9,853	0
1955	6	5,800	1,933	3,867	2,333
1956	12	22,410	7,470	14,940	5,867
1957	19	26,470	8,823	17,647	667
1958	3	16,000	5,333	10,667	1,800
1959	2	10,800	3,600	7,200	6,000
1960	4	15,600	5,200	10,400	0
1961	10	11,295	3,765	7,530	0
1962	3	31,000	10,333	20,667	3,667
1963	2	31,500	10,500	21,000	0

Year	Sales	Gross	Commission	Net	Receipted
1964	2	31,500	10,500	21,000	0
1965	1	15,000	5,000	10,000	10,000
1966	1	25,000	8,333	16,667	0

Table 11: The thirty sales at the highest prices among the rows that name a work, as the work books write them. Gross is the price; net is after the third.

Work	Price	Net	Year	Volume
Intermission	25,000	16,667	1966	III
Sun in an Empty Room	25,000	16,667	1964	III
New York Office	16,500	11,000	1963	III
People in the Sun	15,000	10,000	1962	III
A Woman in the Sun	15,000	10,000	1963	III
East River	15,000	10,000	1965	III
Excursion into Philosophy	14,500	9,667	1962	III
« Girlie Show »	11,000	7,333	1961	II
November, Wash. Square	10,000	6,667	1960	III
Sunlight in a Cafeteria	9,000	6,000	1959	III
Hotel Window	7,000	4,667	1957	III
Western Motel	6,500	4,333	1958	III
Charleston Slum	6,500	4,333	1964	III
Morning in a City	6,000	4,000	1949	III
Four Lane Road	6,000	4,000	1957	III
Dauphinée House	5,000	3,333	1956	II
Sunlight on Brownstones	5,000	3,333	1958	III
Tables for Ladies	4,500	3,000	1930	I
Men, Women + Horses	4,500	3,000	1958	II
Sea Watchers	4,500	3,000	1956	III
South Carolina Morning	4,500	3,000	1957	III
Cape Cod Morning	4,050	2,700	1956	III
« Bootleggers »	4,000	2,667	1956	II
Rooms by the Sea	4,000	2,667	1952	III
The Martha McKeen of Wellfleet	3,500	2,333	1954	III
Portrait of Orleans	3,500	2,333	1953	III
First Row Orchestra	3,500	2,333	1954	III
Morning Sun	3,500	2,333	1954	III
Hotel by a Railroad	3,500	2,333	1954	III
City Sunlight	3,500	2,333	1955	III